

Auto Claims Document Classification - Model Comparison

18 short documents + 12 length probes | Updated 23 September 2026 | OpenJev run incomplete: anonymous demo quota

Jev

Hosted decision model / OpenRouter

Initial documents

10/10

Confusing category cases

8/8

Long probes / full input

12/12

BOUNDARY CASE LATENCY

0.25 s

median / document

OBSERVED INFERENCE COST

\$0.004184

30 total calls; 0.418 cents.
All tests stayed within the \$0.05 cap.

Laya

English checkpoint / local CPU

Initial documents

7/10

Confusing category cases

4/8

Long probes / chunked

12/12

Prefix-only control: 4/12 correct

BOUNDARY CASE LATENCY

0.28 s

median / document

OBSERVED INFERENCE COST

Local compute

No API fee. CPU, memory and setup costs
were not monetized in this pilot.

GPT-5.6 Luna

General LLM / Codex login

Initial documents

10/10

Confusing category cases

8/8

Long probes / full input

12/12

BOUNDARY CASE LATENCY

3.40 s

median / document

API-EQUIVALENT COST ESTIMATE

~\$0.049 estimated

API-equivalent estimate, not a Codex bill.
21 calls / 30 docs; includes overhead. [S5]

OpenJev FP8

Community demo / partial run

Initial documents: 2/2 completed

8 of 10 not run

Boundary cases: not run

Long probes: not run

Quota blocked the third document

PERFORMANCE NOT ESTABLISHED

Only two scored requests completed.
No meaningful speed comparison.

OBSERVED INFERENCE COST

No paid API calls

Free quota exhausted; no paid fallback.
Demo differs from official helper. [S6]

AT YOUR 5k DOCUMENT LENGTH Jev + Luna: 3/3 each. Both received full text. Evidence was tested at start, middle and end.	Laya chunked: 3/3 correct. 324 document-token allowance in this setup; 18 chunks covered each 5k input.	5k observed response times Jev 0.30-0.46 s; Luna 3.22-3.46 s; Laya chunks 7.73-7.74 s.
---	---	--

PUBLISHED CAPABILITIES / DISTINCT FROM TESTED LIMITS

Jev Text only: strings, JSON, text arrays. No image/audio/video. [S1] Direct: 64k total request; state + longest question 32k. OpenRouter: 32k. Service limits, not disclosed architecture. [S1,S2]	Laya Text / JSON-as-text. No documented native image/audio/video path. [S3] Default: 512 total; 324 document tokens here. Backbone: 8,192 positions, not validated Laya capacity. Multilingual: 1,024 default, up to 8,192 encoder. [S3,S4]	GPT-5.6 Luna Text + image input; text output. No audio/video. [S5] Published: 1,050,000 context; 128,000 max output. Codex route may differ. Only text to ~16k document tokens tested. [S5]	OpenJev Official: text, JSON/DOM and one screenshot; 16,384-token prompt limit; 52 choices/pass. MLX: text only. [S6] Demo: text, up to 26 labels; FP8 dequantized to BF16, raw softmax probabilities. CC BY-NC weights; commercial restriction.
---	---	---	--

Confidence is not accuracy. Laya returned native confidence 0.9999 and 0.9943 on two wrong boundary labels. Scores were not calibrated on representative claims.

Limits of the comparison. Long probes reuse one police report in repeated filler at 1k, 4k, 5k and 16k Laya tokens. Synthetic results do not estimate production accuracy.

Text-only evaluation; OCR/images not tested. Timing includes each serving stack: Jev network, Laya CPU, Luna CLI/Codex overhead. Not an intrinsic speed or dollar-cost ranking.
Luna estimate: 232,898 uncached + 99,840 cached input; 471 output. Rates / 1M: \$0.20 / \$0.02 / \$1.20. No-cache estimate: \$0.067. [S5]

PRIMARY EVIDENCE / ORIGINAL SOURCES 22 SEPTEMBER; OPENJEV 23 SEPTEMBER 2026

S1 <https://docs.typesafe.ai/models>

S2 <https://openrouter.ai/typesafe/jev-1.13>

S3 <https://huggingface.co/convaiinnovations/laya>

S4 <https://huggingface.co/answerdotai/ModernBERT-large/blob/main/config.json>

S5 <https://developers.openai.com/api/docs/models/gpt-5.6-luna>

S6 <https://huggingface.co/openjev/openjev>

Demo: <https://huggingface.co/spaces/chanoian/openjev-mlx-demo>

OpenJev: pilot excluded; 2 scored successes, third call quota-rejected. Different prompt/serving stack. No calibrated native confidence returned. Full benchmark remains incomplete.