

Auto Claims Document Classification - Model Comparison

18 short documents + 12 length probes | 22 September 2026 | No additional inference for this visual

Jev Hosted decision model / OpenRouter Initial documents 10/10 Confusing category cases 8/8 Long probes / full input 12/12 BOUNDARY CASE LATENCY 0.25 s median / document OBSERVED INFERENCE COST \$0.004184 30 total calls; 0.418 cents. All tests stayed within the \$0.05 cap.	Laya English checkpoint / local CPU Initial documents 7/10 Confusing category cases 4/8 Long probes / default prefix 4/12 With chunking: 12/12 correct BOUNDARY CASE LATENCY 0.28 s median / document OBSERVED INFERENCE COST Local compute No API fee. CPU, memory and setup costs were not monetized in this pilot.	GPT-5.6 Luna General LLM / Codex login Initial documents 10/10 Confusing category cases 8/8 Long probes / full input 12/12 BOUNDARY CASE LATENCY 3.40 s median / document OBSERVED INFERENCE COST Codex allowance No OpenRouter use. Subscription usage was not converted to a dollar cost.
--	--	--

AT YOUR 5k DOCUMENT LENGTH Jev + Luna: 3/3 each. Both received full text. Evidence was tested at start, middle and end.	Laya: 1/3 prefix; 3/3 chunked. 324 document-token allowance in this setup; 18 chunks covered each 5k input.	5k observed response times Jev 0.30-0.46 s; Luna 3.22-3.46 s; Laya chunks 7.73-7.74 s.
---	---	--

PUBLISHED CAPABILITIES / DISTINCT FROM TESTED LIMITS

Jev Text only: strings, JSON, text arrays. No image/audio/video. [S1] Direct: 64k total request; state + longest question 32k. OpenRouter: 32k. Service limits, not disclosed architecture. [S1,S2]	Laya Text / JSON-as-text. No documented native image/audio/video path. [S3] Default: 512 total; 324 document tokens here. Backbone: 8,192 positions, not validated Laya capacity. Multilingual: 1,024 default, up to 8,192 encoder. [S3,S4]	GPT-5.6 Luna Text + image input; text output. No audio/video. [S5] Published: 1,050,000 context; 128,000 max output. Codex route may differ. Only text to ~16k document tokens tested. [S5]
---	---	---

Confidence is not accuracy. Laya returned native confidence 0.9999 and 0.9943 on two wrong boundary labels. Scores were not calibrated on representative claims.

Limits of the comparison. Long probes reuse one police report in repeated filler at 1k, 4k, 5k and 16k Laya tokens. Synthetic results do not estimate production accuracy.

Text-only evaluation; OCR/images not tested. Timing includes each serving stack: Jev network, Laya CPU, Luna CLI/Codex overhead. Not an intrinsic speed or dollar-cost ranking.
Sources: retained results + report references for published limits. Tokenizers differ. Boundary cases were designed after initial errors; labels frozen before inference.

PRIMARY EVIDENCE / VERIFIED 22 SEPTEMBER 2026

- S1 <https://docs.typesafe.ai/models>
- S2 <https://openrouter.ai/typesafe/jev-1.13>
- S3 <https://huggingface.co/convaiinnovations/laya>
- S4 <https://huggingface.co/answerdotai/ModernBERT-large/blob/main/config.json>
- S5 <https://developers.openai.com/api/docs/models/gpt-5.6-luna>