

Three models. Measured results.

18 short documents + 12 length probes | 22 September 2026 | No additional inference for this visual

Jev

Hosted decision model / OpenRouter

Initial documents	10/10
Confusing category cases	8/8
Long probes / full input	12/12

BOUNDARY CASE LATENCY

0.25 s median / document

OBSERVED INFERENCE COST

\$0.004184

30 total calls; 0.418 cents.
All tests stayed within the \$0.05 cap.

Laya

English checkpoint / local CPU

Initial documents	7/10
Confusing category cases	4/8
Long probes / default prefix	4/12

With chunking: 12/12 correct

BOUNDARY CASE LATENCY

0.28 s median / document

OBSERVED INFERENCE COST

Local compute

No API fee. CPU, memory and setup costs were not monetized in this pilot.

GPT-5.6 Luna

General LLM / Codex login

Initial documents	10/10
Confusing category cases	8/8
Long probes / full input	12/12

BOUNDARY CASE LATENCY

3.40 s median / document

OBSERVED INFERENCE COST

Codex allowance

No OpenRouter use. Subscription usage was not converted to a dollar cost.

AT YOUR 5k DOCUMENT LENGTH

Jev + Luna: 3/3 each. Both received full text. Evidence was tested at start, middle and end.

Laya: 1/3 prefix; 3/3 chunked.

324 document-token allowance in this setup; 18 chunks covered each 5k input.

5k observed response times

Jev 0.30-0.46 s; Luna 3.22-3.46 s;
Laya chunks 7.73-7.74 s.

PUBLISHED CAPABILITIES / DISTINCT FROM TESTED LIMITS

Jev

Text only: strings, JSON, text arrays. No image/audio/video. [S1]
Direct: 64k total request; state + longest question 32k. **OpenRouter:** 32k. Service limits, not disclosed architecture. [S1,S2]

Laya

Text / JSON-as-text. No documented native image/audio/video path. [S3]
Default: 512 total; 324 document tokens here. Backbone: 8,192 positions, not validated Laya capacity. Multilingual: 1,024 default, up to 8,192 encoder. [S3,S4]

GPT-5.6 Luna

Text + image input; text output. No audio/video. [S5]
Published: 1,050,000 context; 128,000 max output. Codex route may differ. Only text to ~16k document tokens tested. [S5]

Confidence is not accuracy. Laya returned native confidence 0.9999 and 0.9943 on two wrong boundary labels. Scores were not calibrated on representative claims.

Limits of the comparison. Long probes reuse one police report in repeated filler at 1k, 4k, 5k and 16k Laya tokens. Synthetic results do not estimate production accuracy.

Text-only evaluation; OCR/images not tested. Timing includes each serving stack: Jev network, Laya CPU, Luna CLI/Codex overhead. Not an intrinsic speed or dollar-cost ranking.
Sources: retained results + report references for published limits. Tokenizers differ. Boundary cases were designed after initial errors; labels frozen before inference.

PRIMARY EVIDENCE / VERIFIED 22 SEPTEMBER 2026

S1 <https://docs.typesafe.ai/models>

S2 <https://openrouter.ai/typesafe/jev-1.13>

S3 <https://huggingface.co/convaiinnovations/laya>

S4 <https://huggingface.co/answerdotai/ModernBERT-large/blob/main/config.json>

S5 <https://developers.openai.com/api/docs/models/gpt-5.6-luna>