

Three models. Measured results.

18 short documents + 12 length probes | 22 September 2026 | No additional inference for this visual

Jev

Hosted decision model / OpenRouter

Initial documents	10/10
Confusing category cases	8/8
Long probes / full input	12/12

BOUNDARY CASE LATENCY

0.25 s median / document

OBSERVED INFERENCE COST

\$0.004184

30 total calls; 0.418 cents.
All tests stayed within the \$0.05 cap.

Laya

English checkpoint / local CPU

Initial documents	7/10
Confusing category cases	4/8
Long probes / default prefix	4/12

With chunking: 12/12 correct

BOUNDARY CASE LATENCY

0.28 s median / document

OBSERVED INFERENCE COST

Local compute

No API fee. CPU, memory and setup costs were not monetized in this pilot.

GPT-5.6 Luna

General LLM / Codex login

Initial documents	10/10
Confusing category cases	8/8
Long probes / full input	12/12

BOUNDARY CASE LATENCY

3.40 s median / document

OBSERVED INFERENCE COST

Codex allowance

No OpenRouter use. Subscription usage was not converted to a dollar cost.

AT YOUR 5k DOCUMENT LENGTH

Jev + Luna: 3/3 each. Both received full text. Evidence was tested at start, middle and end.

Laya: 1/3 prefix; 3/3 chunked.
324 document-token allowance in this setup; 18 chunks covered each 5k input.

5k observed response times

Jev 0.30-0.46 s; Luna 3.22-3.46 s;
Laya chunks 7.73-7.74 s.

INPUT AND OUTPUT CONSTRAINTS

Jev

Text documents tested. OpenRouter lists 32k context tokens; reserve room for question and labels. Returns choice, class probabilities and confidence.

Laya

Text documents tested. Default 512 total tokens left 324 for the document. Later text is silently clipped. Chunking adds compute; pooled confidence is not calibrated.

GPT-5.6 Luna

Text documents tested. Published model context: 1.05M tokens; this Codex route was tested only to ~16k document tokens. Returned JSON label; no native score in this run.

Confidence is not accuracy. Laya returned native confidence 0.9999 and 0.9943 on two wrong boundary labels. Scores were not calibrated on representative claims.

Limits of the comparison. Long probes reuse one police report in repeated filler at 1k, 4k, 5k and 16k Laya tokens. Synthetic results do not estimate production accuracy.